

PARALLEL COMPUTATION OF CLOSENESS CENTRALITY USING APSPS

CS 5990 Computational Social Systems

Team Members:

Aumkareshwar, Jeremy Anunwah, Yurii Lebid, Bill Kim

Closeness Centrality

INTRODUCTION

Objective: Efficiently compute closeness centrality for nodes in large graphs using parallel computing.

Motivation: Closeness centrality is vital for analyzing social networks, but its computation is expensive for large graphs. Parallel computation reduces runtime by distributing workload across processors.

PROBLEM DESCRIPTION

Compute closeness centrality for a graph:

$$C(u) = \frac{n - 1}{\sum_{v=1}^{n-1} d(v, u)}$$

Input:

- Social network datasets (e.g., Facebook, Twitter).

Output:

- Centrality values for all nodes.
- Top 5 nodes with highest centrality.
- Average centrality value.

Algorithm Overview

Key Steps:

- Load and distribute graph data across processors.
- Use Dijkstra's algorithm to compute shortest paths for Subset of nodes.
- Compute closeness centrality locally on each processor.
- Gather and combine results at the root processor.
- Output the centrality measures and performance results.

Pseudocode

Main Process:

- Divide nodes into P subsets.
- Broadcast graph to all processors.
- Compute shortest paths using Dijkstra's algorithm.
- Calculate closeness centrality for assigned nodes.
- Gather results and compute top nodes & average centrality.

Dijkstra's Algorithm:

- Initialize distances as infinity.
- Use priority queue to compute shortest paths.
- Return distances to all other nodes.

Data Structures

Graph Representation:

- Adjacency list (efficient memory usage).

Distance Storage:

- Dictionary for shortest paths.
- Per-processor subset reduces memory usage.

Priority Queue:

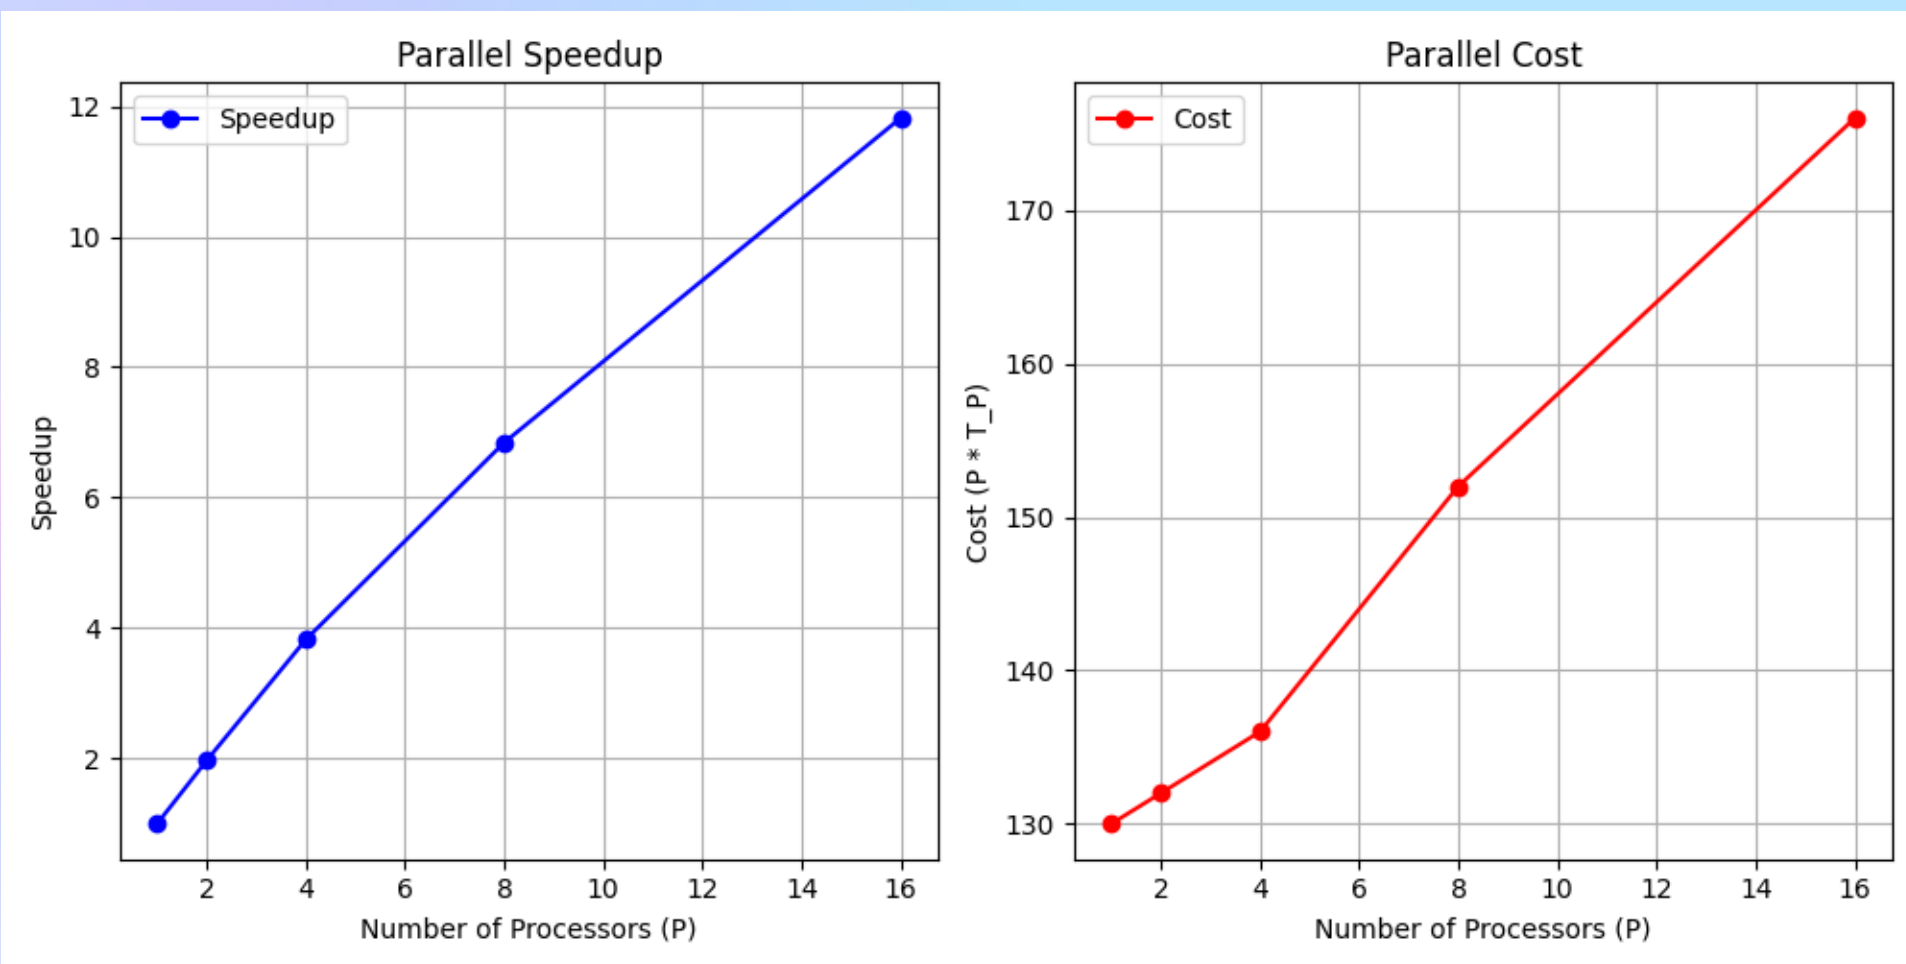
- Implements Dijkstra's algorithm efficiently.

Performance Study Results

Runtime Results:

Process ors (PP)	Runtime (T_{PTP})	Speedup (T_1/TP)	Cost ($P \cdot TP$)
1	130 Seconds	1.0	130
2	66 Seconds	1.9	132
4	34 Seconds	3.8	136
8	19 Seconds	6.8	152
16	11 Seconds	11.8	176

Performance Plots Facebook Data



Performance Plots for Twitter data

Complexity Analysis

Time Complexity:

- Single Processor: $O(n \cdot (n+m) \log n)$
- Parallel Version: $O(n/P \cdot (n+m) \log n) + O(n/P)$

Space Complexity:

- Adjacency list: $O(n+m)$
- Per-processor storage: $O(n/P)$

Strengths and Limitations

Strengths:

- Efficient for sparse graphs.
- Scalable for small to moderate P .

Limitations:

- Diminishing returns at high processor counts due to communication overhead.
- Load imbalance for graphs with non-uniform node degrees.

Output for Facebook and Twitter Data

```
asrinivas1@commander:~  
(base) [asrinivas1@cn12 ~]$  
(base) [asrinivas1@cn12 ~]$  
(base) [asrinivas1@cn12 ~]$  
(base) [asrinivas1@cn12 ~]$  
(base) [asrinivas1@cn12 ~]$  
(base) [asrinivas1@cn12 ~]$  
(base) [asrinivas1@cn12 ~]$  
(base) [asrinivas1@cn12 ~]$  
(base) [asrinivas1@cn12 ~]$  
(base) [asrinivas1@cn12 ~]$ nano dijk.py  
(base) [asrinivas1@cn12 ~]$ ls  
asps.py      dijk.py      output.txt   sn-hpc.py  
centrality.py facebook_combined.txt quakers_edgelist.csv  
close.py     miniconda3   quakers_nodelist.csv  
(base) [asrinivas1@cn12 ~]$ mpiexec -n 8 python dijk.py  
Loading graph...  
Top 5 nodes with highest closeness centrality:  
Node 107: 0.45969945355191255  
Node 58: 0.3974018305284913  
Node 428: 0.3948371956585509  
Node 563: 0.3939127889961955  
Node 1684: 0.39360561458231796  
Average Closeness Centrality: 0.2761677635668376  
(base) [asrinivas1@cn12 ~]$
```

Facebook Data

References

- SNAP Datasets: [Facebook Dataset](#), [Twitter Dataset](#).
- MPI Documentation: [mpi4py](#).

Tasks

- Jeremy Anunwah: presentation, simulation experiments, contributed to research, algorithm design, MPI communication.
- Yurii Lebid: Presentation, Prepared datasets, optimized algorithm design, and conducted performance analysis.
- Bill Kim: optimization, debugging, and presentation preparation.
- Aumkareshwar: presentation, algorithm implementations, HPC setup, testing, and generating performance plots.



THANK YOU